

Simple Neural Network for OR Gate Classification

1. Project Overview

This project implements a very simple **single-layer neural network** (a **perceptron**) using pure Python, without relying on any machine learning libraries like TensorFlow, Keras, or PyTorch.

The goal is to:

- Train the network to model the behavior of a **logical OR gate**.
- Track and save the **loss values during training** into a .csv file.
- Use the saved loss values for further **visualization and analysis** in a Jupyter Notebook.

2. Neural Network Details

- **Inputs:** Two binary features (0 or 1) representing logical inputs.
- **Output:** A single output (0 or 1) predicted using a sigmoid activation function.
- **Architecture:**
 - 2 input nodes.
 - 1 output node (no hidden layers).
 - Trainable parameters: weights w_1 , w_2 and bias b .
- **Activation Function:**
 - $\text{sigmoid}(x) = 1 / (1 + \exp(-x))$
 - Outputs values between 0 and 1, useful for binary classification.

3. Training Process

- **Loss Function:**
Mean Squared Error (MSE) between predicted and actual outputs:

$$\text{MSE} = \overset{\text{Mean}}{\frac{1}{n} \sum_{i=1}^n} \left(\overset{\text{Error}}{Y_i - \hat{Y}_i} \right) \overset{\text{Squared}}{^2}$$

- **Optimization Algorithm:**

Manual implementation of gradient descent by **finite difference approximation**:

- Approximates partial derivatives by perturbing each parameter slightly ($\epsilon = 0.01$) and calculating change in loss.

- **Hyperparameters:**

- Learning rate (η): 0.1
- Number of epochs: 5000
- Epsilon for finite difference: 0.01

- **Training Dataset:**

Models the logical OR function:

x1	x2	y (output)
0	0	0
1	0	1
0	1	1
1	1	1

- **Outputs After Training:**

- Final values of weights w_1 , w_2 , and bias b .
- Printed predictions for all input combinations.
- Loss values saved to a CSV file (loss_values_gates.csv).

4. Jupyter Notebook Analysis (graphs.ipynb)

The Jupyter notebook:

- Loads the saved loss_values_gates.csv.
- Visualizes the loss values over epochs.
- Helps in observing:
 - How fast the model converges.
 - Stability of the training process.

Main steps inside the notebook:

- Load CSV using pandas.

- Plot loss values using matplotlib.

5. Important Code Highlights

`sigmoid(x)`

Calculates the sigmoid activation.

`loss_function(w1, w2, b, train_data)`

Computes the mean squared error (MSE) between predicted and true outputs for given weights and bias.

`train(train_data)`

Main training function:

- Initializes weights and bias randomly.
- Performs manual gradient descent using finite difference approximation.
- Updates parameters and saves loss values for each epoch.

`graphs.ipynb`

- Reads loss values.
- Plots a line graph of loss versus number of epochs.

6. Outputs

At the end of training, I got this console output:

```
initial loss: [0.13294505]
```

```
-----
```

```
W1 = [4.53749557], W2 = [4.53528661], bias = [-2.00368412], Loss = [0.00623779]
```

```
-----
```


```
0 | 0 = [0.11881666]
```

```
0 | 1 = [0.92632779]
```

```
1 | 0 = [0.9264784]
```

```
1 | 1 = [0.99914972]
```

Then, in the Jupyter Notebook, the loss values recorded during training are plotted.



Conclusion from the Graph:

- The loss started relatively **high** at the beginning of training.
- It then **decreased rapidly** during the early epochs.
- After 5000 epochs, the loss **stabilized** at a very low value, close to **zero**, indicating that the model successfully **learned the logical OR gate**.
- The curve shows a **smooth descent**, suggesting that the learning rate was well-chosen and the model converged **efficiently** without oscillations.